


An introduction to tinc

Hamburg, July 28, 2014



An introduction
to tinc

Guus Sliepen

Introduction

Communicating
over the Internet

The problem of
NAT

The problem of
MTU

Other problems

Future plans for
tinc

Authentication and
authorization

The end

Guus Sliepen

guus@tinc-vpn.org

Before September 1997:

- Bad support for network virtualization
- No IPsec
- PPP module:
 - You have to speak the PPP protocol
 - Or run pppd through something else (like SSH)
 - Very poor performance
- CIPE module
 - out-of-tree
 - Does everything inside the kernel
 - Very inflexible, hard to maintain

September 1997, Linux 2.1.53 is release.

- Introduces the Ethertap module
- ```
modprobe ethertap
ifconfig tap0 ...
fd = open("/dev/tap0", O_RDWR);
```
- Several userspace projects start here.

#### Introduction

Communicating  
over the Internet

The problem of  
NAT

The problem of  
MTU

Other problems

Future plans for  
tinc

Authentication and  
authorization

The end

The very beginning of tinc: "tapencaps"

- Read packets from tap device, send them to TCP socket.
- Read packets from TCP socket, send them to tap device.
- Rest is left to the user.

That's nice but what can you do with that?

- Set up a private network to a friend!

This was started by me.

The next stage: "vpnd"

- Turn tapencaps into a proper UNIX daemon.
- And make releases.
- And put source code into CVS.
- Slap the GPL license on it.

*In short...*

- Become an Open Source project.

This work was mostly done by Ivo Timmermans.

We used it for our own friend network.

- Everyone wants to be able to connect to each other.
- Try a hub-and-spoke setup.
- One vpnd per spoke.
- Hub gets a lot of traffic.
- Falls apart everytime the hub goes down.
- Soon we reach the limit of 16 Ethertap devices.

That's not good. What to do?

December 1998, project renamed "tinc".

January 1999, support for multiple connections:

- One daemon can connect to multiple other daemons
- One daemon still uses only one Ethertap device
- Tinc reads the IP header, routes packets.

This set tinc apart from other VPN projects.

#### Introduction

Communicating  
over the Internet

The problem of  
NAT

The problem of  
MTU

Other problems

Future plans for  
tinc

Authentication and  
authorization

The end

## *Fast-forward* - Current features:

- Connects multiple sites together
- Can act as router (layer 3) or switch (layer 2)
- Full support for IPv6
- No central server
- You configure some endpoints, tinc will do the rest

## Modus operandi:

- Metadata exchanges via TCP
- VPN packets directly via UDP
- Fall back to TCP if UDP is not possible




## The competition:

- CIPE<sup>†</sup>
- VTun<sup>†</sup>
- IPsec
- OpenVPN
- Hamachi

## But also:


- GVPE
- CloudVPN
- SocialVPN
- n2n
- VDE

Network before VPN is configured:




Blue cloud: the Internet  
Black circles: VPN nodes

Initial connections configured by user:




Green lines: initial connections

Full mesh created by tinc:



Black lines: UDP tunnels

Reality is not so nice:



Red arcs: NAT

Red line: ISP blocking traffic

Dotted lines: failed connections


## The problem of NAT:

- Source address and port change
- Incoming connections blocked

## Solutions:


- Routing via non-NAT node (not efficient)
- Port forwarding (not always possible, manual work)
- UPnP (needs router support, complex)
- STUN/ICE (not always possible, complex)

## Two nodes behind NAT:



Both nodes can talk to a third node.  
NAT changes ports, nodes cannot talk to each other.

## STUN in action:




Third node tells other nodes about their addresses and ports.  
Nodes connect to each other using this information.



## The problem of packet fragmentation:

- MTU inside tunnel smaller than outside
- Outer layer fragments bad for performance
- Some firewalls/ISPs block fragments




## Solution:

- Determine path MTU between nodes
- Generate ICMP Fragmentation Needed packets
- Should work for all IPv4/IPv6 traffic
- Fall back to TCP for other traffic

## The problem of firewalls/ISPs blocking ICMP:

- ICMP Fragmentation Needed does not work!
- Happens when network traffic leaves the VPN (for example, when having default gateway on VPN)



## Solution:

- Clamp MSS field in TCP packets to path MTU
- Works only for TCP

## Other problems:

- Frequently changing IP addresses
  - use dyndns
  - cache & forward known addresses between nodes
- Only allowing certain ports, like HTTP
  - tunnel over ICMP/DNS/HTTPS
- ISPs dropping/delaying small UDP packets
  - because they think it's VoIP!
  - severely slows down TCP streams inside tunnel

## Version 1.1:

- Control a running tinc daemon (CLI, GUI)
- Greatly simplify setting up tinc
- Improve security (SPTPS)

## Version 2.0:

- Get rid of the legacy protocol
- Decentralized authorization framework (next slides)



### An introduction to tinc

Guus Sliepen

[Introduction](#)

[Communicating over the Internet](#)

[The problem of NAT](#)

[The problem of MTU](#)

[Other problems](#)

[Future plans for tinc](#)

[Authentication and authorization](#)

[The end](#)

# Authentication and authorization

- Authentication = proving who you are
- Authorization = proving you are allowed to do something



## Two well known (mostly authentication) methods:

- X.509 certificates
  - centralized approach
  - focused on identities (LDAP like), and URLs
- OpenPGP keys
  - decentralized approach
  - focused on email addresses

Introduction

Communicating  
over the Internet

The problem of  
NAT

The problem of  
MTU

Other problems

Future plans for  
tinc

Authentication and  
authorization

The end

We need OpenPGP-like features:

- Completely decentralized
- Web of trust

We need more than OpenPGP can offer:

- Authorise anything, not just email
- Everyone can add/remove authorizations
- Negative authorization: forbid things
- Group decisions

## libfides: lightweight p2p authorization framework

- Create and maintain repository of many certificates  
*"X said at time  $T$  that  $Y$  is allowed to do  $Z$ "*
- Newer certificates overrule older ones  
*"X said at time  $T+1$  that  $Y$  is not allowed to do  $Z$ "*
- Make it simple to query the repository  
*"Is  $Y$  allowed to do  $Z$ ?"*
- Fast & easy synchronization of repositories
- Application does not need to know about crypto
- Libfides itself uses only ECDSA primitive
- Still in alpha stage.

### An introduction to tinc

Guus Sliepen

Introduction

Communicating over the Internet

The problem of NAT

The problem of MTU

Other problems

Future plans for tinc

Authentication and authorization

The end

## Conclusions:

- Tinc has come a long way since 1997.
- The Internet eats your packets.
- Lots of techniques necessary to work around it.
- Distributed authorization is a challenge.

That's it.

- Questions?

Visit the website:

<http://tinc-vpn.org/>

